DISTRIBUTION-SHIFTED SYNTHETIC MULTITRACK PRETRAINING FOR LARGE-VOCABULARY AUDIO CHORD RECOGNITION

Sharabh Ojha
University of California, Los Angeles
shojha@ucla.edu

ABSTRACT

Audio chord recognition (ACR) models trained on human-annotated corpora face a persistent triadic performance plateau: canonical benchmarks concentrate over 80% of duration in major and minor triads, leaving rare triads (`sus2`, `sus4`, `dim`, `aug`) and upper extensions (7th–13th) acoustically sample-starved. Consequently, neural classifiers underperform on minority harmonies, and we hypothesize that inverse-frequency loss reweighting alone cannot compensate for the scarcity of diverse acoustic exemplars. We propose a Sim-to-Real pretraining framework driven by a probabilistic composition generator that synthesizes $5{,}000$ multitrack audio tracks (460.3 h) using rare-chord quota scheduling, physically constrained multi-instrument voicings, and stochastic random-mix acoustic augmentation. We pretrain ChordFormer—a 4-block Conformer encoder with six decomposed output heads—on synthetic multitracks and finetune on real audio via inverted discriminative learning rates ($\mathbf{s}_{\text{lr}} = [0.2, 0.5, 1.0]$) and floor curriculum replay ($\rho_{\min} = 0.10$). Across all $11{,}436{,}357$ frames of the $1{,}217$-song benchmark, our curriculum achieves substantial recall gains over the real-only baseline on underrepresented harmonies: boosting rare triad recall from 41.67% to 48.31% ($+6.64$ pp; F_1 from 53.11% to 60.94%), 9th extensions from 35.33% to 43.58% ($+8.25$ pp), natural 9ths by $+25.16$ pp ($31.64\% \rightarrow 56.80\%$), altered `#9` by $79\times$ ($0.07\% \rightarrow 5.53\%$), and flat 13ths by $7.8\times$ ($0.84\% \rightarrow 6.58\%$). While duration-weighted MIREX WCSR shifts from 81.49% to 79.59% due to metric bias toward dominant triads, cross-fold variance contracts across all ten evaluation rules.

1. INTRODUCTION

Automatic Chord Recognition (ACR) underpins computational musicology, structural segmentation, and content-based audio indexing [1]. Although ACR architectures have progressed from handcrafted chroma templates and hidden Markov models to learned deep chroma extractors [2], convolutional-recurrent networks [3], and self-attention Transformers [4, 5], performance on complex chords remains bounded: classifiers recognize major and minor triads far better than four-note tetrads and extended harmonies [1] (e.g., 77.6% MajMin versus 58.8% Tetrads weighted recall for a K-stream HMM on the canonical $1{,}217$-song benchmark [6]).

This ceiling stems not from insufficient model capacity, but from the severe Zipfian skew of real-world musical corpora: in the $1{,}217$-song benchmark, major (65.75%) and minor (22.16%) triad qualities (Head 0, including extended chords) occupy 87.91% of non-silent duration (73.7 h, $11{,}436{,}357$ frames). Rare triads (`sus4`, `dim`, `sus2`, `aug`) collectively span under 3.5% of duration, while altered upper extensions (`#11` at 0.13%, `b9` at 0.13%, `b13` at 0.08%) inhabit a sparse tail. Moreover, human inter-annotator disagreement exceeds 40% on seventh chords, tetrads, and their inversions [7], compounding sample scarcity with label noise.

Recent architectures address large-vocabulary recognition via structured multi-task losses [8], 6-head chord decomposition [3, 9], and contrastive pseudo-labeling [10]. Specifically, Jiang et al. [3] and Akram et al. [9] factorize chords into six categorical heads (root+triad, bass, and $7\text{th}, 9\text{th}, 11\text{th}, 13\text{th}$ degrees) trained under inverse-frequency reweighted cross-entropy. We hypothesize that loss reweighting alone cannot overcome a scarcity of physical acoustic exemplars: when an altered degree occurs in only a few dozen frames across a corpus, scaling its gradient penalty may amplify gradient variance and overfit to the instrumentation and room acoustics of those few tracks. Directly addressing the long tail motivates intervening on the training audio distribution.

We propose a Sim-to-Real pretraining framework driven by procedurally synthesized multitrack audio, contributing: (1) **Probabilistic Composition Generator**, an open-source engine sampling progressions along empirical transitions curated via the ChoCo aggregator [11] under deterministic quota scheduling (upweighting rare triads up to $15.6\times$ and altered extensions up to $12.3\times$), enforcing physical register bounds and 100% extension sounding across four stems with random-mix augmentation to render $5{,}000$ tracks (460.3 h) and frame-exact Harte labels [12]; (2) **Two-Stage Sim-to-Real Curriculum**, a transfer regimen for ChordFormer [9] combining synthetic pretraining with real-audio finetuning via inverted discriminative learning rates ($\mathbf{s}_{\text{lr}} = [0.2, 0.5, 1.0]$) and floor curriculum replay ($\rho(\eta) = 0.150 \rightarrow 0.100$); and (3) **Multi-Head and WCSR Benchmark Analysis**, a 5-fold cross-validated evaluation over all $11{,}436{,}357$ frames demonstrating substantial rare-chord recall gains over the real-only baseline ($+6.64$ pp on rare triads, $+25.16$ pp on natural 9ths, $79\times$ on `#9`, $7.8\times$ on `b13`) alongside an analysis of precision-recall trade-offs, duration-weighted WCSR bias, and poly-

phonic overtone masking.

2. RELATED WORK

2.1 Large-Vocabulary and Structured Chord Recognition

Early ACR systems restricted classification to 24 major/minor triads using chroma vectors and hidden Markov models [1, 2]. To scale to larger vocabularies without combinatorial explosion, McFee and Bello [8] formulated a structured multi-task loss over 170 classes (14 qualities), regularizing predictions via root, bass, and pitch-class targets; however, their taxonomy maps all upper extensions (9th, 11th, 13th) onto basic triads or tetrads.

Jiang et al. [3] preserved upper tertian degrees by decomposing chords into six categorical heads (Root+Triad, Bass, 7th, 9th, 11th, 13th) trained with CRNNs, inverse-frequency class reweighting, and CRF decoding. Self-attention architectures further improved temporal context modeling: the Bidirectional Transformer for Chords (BTC) [4] replaced recurrent layers with bidirectional self-attention, and Kim and Park [13] formulated segment-level sequence-to-sequence decoding to suppress frame over-segmentation on non-triad harmonies.

Akram et al. [9] proposed ChordFormer, pairing Macaron-style Conformer blocks [5] with the 6-head decomposition of Jiang et al. [3] to capture local harmonic partials via depthwise convolution and global phrase structure via multi-head attention. Li et al. [10] applied contrastive learning and Noisy Student pseudo-labeling to unannotated audio across 169 classes. However, teacher-student pseudo-labeling suffers from confirmation bias on the harmonic tail: because teacher networks exhibit near-zero recall on rare alterations, pseudo-labels reinforce majority diatonic priors. Our approach instead supplies noise-free, symbolically verified ground-truth audio across the full extension vocabulary.

2.2 Synthetic Audio Datasets in MIR

Synthetic multitrack rendering offers a scalable remedy for annotation scarcity in MIR. In source separation, Slakh2100 [14] rendered 2,100 MIDI arrangements through virtual instrument libraries to provide isolated stems. For chord recognition, Ostermann et al. [15] developed Artificial Audio Multitracks (AAM), comprising 3,000 sampler-rendered tracks generated via algorithmic composition in JFugue.

Majchrzak and Mańdziuk [16] benchmarked AAM on BTC and Harmony Transformer models under a 25-class major/minor vocabulary, showing that combining AAM with real audio improves pop WCSR (+2.68 pp on Billboard), whereas training solely on synthetic tracks degrades accuracy on human-composed pop and classical recordings. Because AAM samples progressions exclusively from unextended diatonic triads (I–vi) [15], it replicates the triadic skew of real corpora and precludes large-vocabulary transfer. In contrast, our generator acts as an explicit distributional intervention: upweighting rare triads and altered extensions via quota scheduling, enforcing physical voicings so labeled extensions sound acoustically, and applying random-mix augmentation to bridge the SoundFont-to-studio gap.

3. PROBABILISTIC COMPOSITION GENERATOR

3.1 Distribution-Shifted Harmonic Sampling

Whereas AAM [15] renders 3,000 tracks from random diatonic triads (I–vi), omitting vii°, suspended/augmented triads, and 7th–13th extensions [16], we design `target_corpus_gen.py` as a distributional intervention over 5,000 tracks (3,250 pop/rock, $r_{\text{pop}} = 0.65$; 1,750 jazz, 0.35; 500,000 events, 100/track) along transitions extracted via ChoCo [11] from McGill Billboard [17] and Isophonics [18] (1,114 songs, 139,078 events) plus the Weimar Jazz Database [19] and iReal Pro [1] (2,477 charts, 161,022 events).

We parameterize each event C_t by extending the 6-head space of Jiang et al. [3] and ChordFormer [9]: whereas both works restrict Head 0 to 6 triad qualities (73 classes, 100 logits)—omitting incomplete triads `5` and `1` in training and collapsing them into `maj` at evaluation—synthesizing standalone `5` and `1` exemplars allows us to expand Head 0 to 8 qualities $q_t \in \{\texttt{maj}, \texttt{min}, \texttt{sus4}, \texttt{sus2}, \texttt{dim}, \texttt{aug}, 5, 1\}$ across roots $r_t \in \mathbb{Z}_{12}$ to form $S_t = (r_t, q_t)$ ($\hat{\mathbf{z}}_{0,t} \in \mathbb{R}^{97}$ with `N`); bass $b_t = (r_t + \Delta_{\text{bass},t}) \bmod 12$ for $\Delta_{\text{bass},t} \in \mathbb{Z}_{12}$ ($\hat{\mathbf{z}}_{1,t} \in \mathbb{R}^{13}$); and $\mathbf{E}_t = (e_{7,t}, e_{9,t}, e_{11,t}, e_{13,t})$ spans $\{\texttt{N}, 7, \texttt{b7}, \texttt{bb7}\}$ ($\hat{\mathbf{z}}_{2,t} \in \mathbb{R}^4$, mapping `dim7` and, following ChordFormer [9], `maj6`/`min6` to `bb7`; Jiang et al. [3] instead map the sixth to the 13th head), $\{\texttt{N}, 9, \texttt{\#9}, \texttt{b9}\}$ ($\hat{\mathbf{z}}_{3,t} \in \mathbb{R}^4$), $\{\texttt{N}, 11, \texttt{\#11}\}$ ($\hat{\mathbf{z}}_{4,t} \in \mathbb{R}^3$), and $\{\texttt{N}, 13, \texttt{b13}\}$ ($\hat{\mathbf{z}}_{5,t} \in \mathbb{R}^3$), totaling 124 logits. Per song (4/4 meter), we sample tonic $k \sim U_{\mathbb{Z}}(0, 11)$ ($S_1 \in \{\texttt{0_major}, \texttt{0_minor}\}$), BPM $\sim U_{\mathbb{Z}}(60, 180)$, and durations $d_t \in \{8, 6, 4, 3, 2, 1.5, 1\}$ (`ww`–`q`) with pop/rock weights $(.02, .03, .60, .05, .20, .05, .05)$ and jazz weights $(.01, .01, .30, .08, .40, .05, .15)$ ($\Delta t_{\text{sec}} = 60 d_t/\text{BPM}$).

In `chord_gen.py`, sequence transitions factorize as:

$$P(C_t \,|\, C_{t-1}, C_{t-2}) = P(S_t \,|\, S_{t-1}, S_{t-2})\, P(b_t \,|\, r_t, b_{t-1}) \cdot P(\mathbf{E}_t \,|\, S_t, b_t, \mathbf{E}_{t-1}). \tag{1}$$

Stage 1 interpolates order-2, order-1, and unigram tables via Witten-Bell weights $\lambda_i = n_i/(n_i + D)$ ($D = 25.0$):

$$P_{\text{final}}(S_t) = \lambda_2 P_{\text{L2}}(S_t \,|\, S_{t-1}, S_{t-2}) + (1 - \lambda_2) \cdot \big[\lambda_1 P_{\text{L1}}(S_t \,|\, S_{t-1}) + (1 - \lambda_1) P_{\text{L0}}(S_t)\big], \tag{2}$$

using tempered $P_{\text{L1}} \propto P_{\text{raw}}^{1/\tau_1}$ ($\tau_{1,\text{pop}} = 2.2$, $\tau_{1,\text{jazz}} = 2.0$), floor $(1 - \epsilon) P_{\text{final}} + \epsilon P_{\text{L0}}$ ($\epsilon = 0.02$), and repeat discount $\delta = 0.20$ on $S_t = S_{t-1}$. Stage 2 applies order-1 bass backoff ($D = 25.0$), adds whole/half-note slash inversions at $p_{\text{inv}} = 0.10$ (0.60 3rd, 0.40 5th), and slices sub-measure passing bass motion. Stage 3 walks a 4-slot prefix trie

[1] `https://www.irealpro.com`

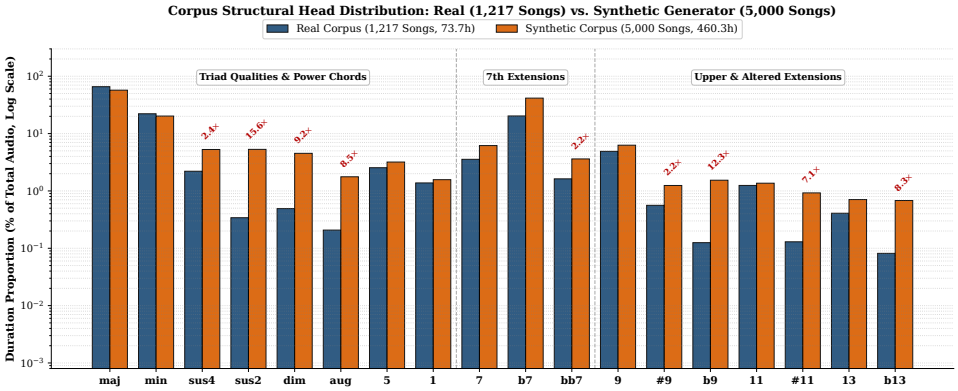


Figure 1. Log-scale duration shares across 18 chord and extension classes in the 1,217-song real benchmark (73.7 h) versus our 5,000-song synthetic corpus (460.3 h, excluding `N`). Head 0 major (65.75%) and minor (22.16%) qualities occupy 87.91% of non-silent real audio, leaving a steep Zipfian tail; quota scheduling expands rare triads (`sus2`: 0.34%→5.31%, 15.6×; `dim`: 9.2×; `aug`: 8.5×; `sus4`: 2.4×) and upper extensions (`b9`: 12.3×; `b13`: 8.3×; `#11`: 7.1×; `#9`: 2.2×).

across state Trie A ($N \geq 200$), triad Trie B ($N \geq 500$), and genre Trie C with $(b_t, \mathbf{E}_{t-1})$ backoff, merges enharmonic collisions (`bb7+13`, `sus4+11`, `sus2+9`, `dim+#11`, `aug+b13`), and injects $p_{\text{NC}} = 0.01$ silence (`N`) solely on unextended root-position `maj`/`min` triads.

We partition tracks ($r_{\text{nat}} = 0.50$) into 2,500 natural and 2,500 targeted songs. For class k with count c_k, quota Q_k ($S_{\text{triad}} = 3.0\times$ on rare triads), event budget N_{ev}, and buffer $B \in \{50, 5{,}000\}$, due positions satisfy:

$$\text{pos}_{\text{due}}(k) = \left\lfloor \frac{(c_k + 1)N_{\text{ev}} + (Q_k + B) - 1}{Q_k + B} \right\rfloor, \quad (3)$$

spaced by $\Delta_{\text{pos}} \geq 2$ along attested 1- or 2-step edges weighted by $P_{\text{dist1}}(s)C_{\text{local}}(s)/(1 + \text{RootCount}(s))$ (3.38×–95.8× empirical upweighting). Versus the 1,217-song corpus (Figure 1), duration shares rise across `sus2` (0.34%→5.31%, 15.55×), `dim` (0.49%→4.53%, 9.20×), `aug` (0.21%→1.77%, 8.50×), `sus4` (2.21%→5.28%, 2.40×), `b7` (20.30%→41.55%), `bb7` (1.62%→3.61%), `b9` (0.13%→1.54%, 12.28×), `b13` (0.08%→0.68%, 8.34×), `#11` (0.13%→0.92%, 7.12×), `#9` (0.56%→1.25%, 2.22×), `9` (4.90%→6.27%), `13` (0.41%→0.71%), and `11` (1.25%→1.36%).

3.2 Symbolic Realization and Voicing Heuristics

Adapting AAM's four-stem JFugue paradigm [15] (`render.py`), we map 4/4 16th-note grids (4 slots/beat) across: `V0` (70% pads, 30% arpeggios), `V1` (bass in $[28, 55]$, E_1–G_3, playing roots, inversions, walking/passing lines, and $p_{\text{jump}} = 0.20$ octave leaps), `V2` (beam-search or random-walk melody masking $\{3, 4\}$ on `sus2`/`sus4`/`1` and restricting `5` to $\{0, 5, 7, 10\}$), and `V9` (Channel 10 drums).

Six policies parameterize `V0` MIDI bounds $[w_{\text{lo}}, w_{\text{hi}}]$ and voice limits $[V_{\text{min}}, V_{\text{max}}]$ [20]: `pop-piano` ($[43, 88]$, $[3, 7]$), `pop-guitar` ($[40, 76]$, $[3, 6]$), `pop-synth` ($[36, 96]$, $[3, 8]$), `jazz-piano` ($[48, 84]$, $[3, 6]$), `jazz-guitar` ($[45, 79]$, $[3, 5]$), and `jazz-synth` ($[40, 91]$, $[3, 7]$), with guitar policies enforcing $\leq$ 4-fret spans (cap $\leq 12/15$) and barré/drop shapes to prevent seventh-chord collapse across fretboard acoustics [21]. Unlike shell voicers that drop extensions, `voicing/select.py` enforces acoustic faithfulness: (1) 100% of 7th–13th extensions sound in piano/synth policies, while guitar policies allow fallback only below Step 8 of an 8-level ladder at $+60.0N_{\text{dropped}}$ penalty ($< 0.5\%$ drop rate); (2) the Distinguishing Chord Tone (13th > 11th > 9th > 7th > 3rd > 5th) avoids upper-root masking via top, isolated, or octave exposure; (3) 3rds never drop and 5ths omit ($p_{\text{omit_5th}} \in [0.10, 0.70]$) only on `maj`/`min`; and (4) `V0` root omission ($p_{\text{omit_root}} \in [0.02, 0.75]$) requires active `V1` bass. After low-interval and cluster pruning, we sample voicings via $P(c) \propto \exp(-(C(c) - \min C)/\tau_{\text{voic}})$ ($\tau_{\text{voic}} \in [2.2, 6.0]$) under:

$$\begin{aligned} C(c) = {} & w_{\text{vl}} d_{\text{vl}}(p_{\text{prev}}, p_c) - B_{\text{plr}} \\ & - 0.5 \min(|p_{\text{prev}} \cap p_c|, 3) + C_{\text{reg}}(c), \end{aligned} \quad (4)$$

combining minimal voice-leading distance d_{vl} [22, 23] (solved via Hungarian assignment; $w_{\text{vl}} \in [0.5, 1.6]$, unmatched penalty $[1.2, 2.5]$), Neo-Riemannian P/L/R bonus $B_{\text{plr}} \in [0.4, 3.0]$ [24, 25], and spacing/drift cost C_{reg}. Five arpeggio profiles (gate 0.65–1.75, velocity 78–96) and `HumanizedMidiRenderer.java` prevent block-onset overfitting via $\sigma_t = 15.0$ ms jitter, pocket offsets $(0, +8, -4, -6, +2$ ms on pads, `V1`, `V2`, arpeggios, `V9`), 7.0 ms/string strums (GM 24–31), CC64 sustain (GM 0–7), and Gaussian velocity noise ($\sigma_v = 8.0$ harmonic; $\sigma_v = 4.0$ atop drum hits $\mathcal{N}(100, 8^2)$ kick/snare/crash, $\mathcal{N}(65, 8^2)$ ghost snare, $\mathcal{N}(80, 12^2)$ hi-hat/ride).

3.3 Synthesis and Random-Mix Augmentation

We synthesize stems via FluidSynth at $f_s = 22{,}050$ Hz (16-bit FLAC), rotating across four GM `.sf2` banks: `FluidR3_GM` (141 MB), `MuseScore_General_Full` (468 MB), `GeneralUser_GS` (31 MB), and `TimGM6mb` (5.7 MB). GM pools span 40 `V0` patches (16 keyboard/organ/mallet, 12 guitar/plucked, 12 synth/pad), 13 weighted `V1` bass patches (GM $33, 34, 32, 35$ weighted $35, 25, 20, 15$), 18 `V2` lead patches ($\geq 25\%$ vocal formants GM 52–54 in pop/rock), and randomized drum kits ($p = 0.60$ crash).

To mitigate acoustic overfitting, we mix stems via `project_condition_corpus.py` using Bernoulli dropout on `V2` and `V9` ($P = 0.70$ each: 49%

four-stem, 21% `cbm`, 21% `cbp`, 9% `cb`) and perturb nominal gains $(0.0, +1.0, -1.0, -1.5)$ dB by $\Delta G_{\text{V0}}, \Delta G_{\text{V1}} \sim U(-3, +3)$ dB, $\Delta G_{\text{V2}} \sim U(-3, +4)$ dB, and $\Delta G_{\text{V9}} \sim U(-4, +4)$ dB. FFmpeg applies `V1` saturation ($x - 0.05x^3$), 120 Hz shelf ($+2.0$ dB), and 3.5 kHz low-pass; `V0` power-chord overdrive and cabinet EQ ($-3.0, +2.5, -6.0$ dB at $0.1, 0.25, 4.5$ kHz); stereo panning ($60/40$ `V0`, $50/50$ `V2`, $40/60$ `V9`); split-band room convolution ($1.2/0.8/0.4$ s, $82\%/18\%$ dry/wet, 150 Hz wet high-pass); detuning $\Delta f_{\text{cents}} \sim U(-25, +25)$ cents; $\pm U(1.5, 3.0)$ dB tilt at 1.5 kHz; console saturation ($x - 0.1x^3$); bus compression ($20/150$ ms); and -18.0 LUFS normalization (-1.0 dBTP ceiling). This forces the Conformer [5, 9] to extract timbre-invariant pitch-class profiles across polyphonic masking and detuned CQT bins. Across all 5,000 tracks (460.3 h), `start_padding_frames=0` anchors sample 0 to $t = 0.000$ s, exporting frame-exact Harte `.lab` labels [12] free of human annotation subjectivity [7].

4. EXPERIMENTAL SETUP

4.1 Audio Preprocessing and ChordFormer Architecture

Audio is downmixed to mono at $f_s = 22{,}050$ Hz [9]. We extract hybrid CQT spectrograms [26] with hop $H = 512$ samples ($\Delta t = 23.22$ ms), window $N_{\text{win}} = 8{,}192$ samples, and $b = 36$ bins/octave (33.3 cents/bin) spanning $F_{\text{raw}} = 288$ bins (8.0 octaves) from $f_{\min} = 23.12$ Hz ($\text{F}\sharp_0$, $A_4 = 440.0$ Hz). Online pitch shifting steps across $[-5, +6]$ semitones on a 3-bin grid, slicing $F = 252$ bins (7.0 octaves, $\text{C}_1 = 32.70$ Hz to $\text{B}_7 = 3{,}951.07$ Hz). Log-amplitude scaling clips to $[-80.0, 0.0]$ dB, and right-zero-padding by 4 bins yields $d_{\text{in}} = 256$ channels ($\mathbf{x} \in \mathbb{R}^{B \times T \times 256}$).

Our `confor_head16.py` adapts ChordFormer [9], stacking $N_L = 4$ Macaron Conformer blocks [5] ($d_{\text{model}} = 256$, $p_{\text{drop}} = 0.10$): (1) half-step FFN_1 ($d_{\text{ffn}} = 1{,}024$); (2) depthwise convolution ($K = 31$ frames ≈ 720 ms, `SAME` pad 15, Pointwise $256 \to 512 \to$ GLU, Depthwise 256, 1-group GroupNorm, SiLU, Pointwise $256 \to 256$, `convolution_first=True`); (3) 16-head MHSA ($h = 16$, $d_k = 16$); (4) half-step FFN_2; and (5) LayerNorm. Unlike the original [9] ($h = 4$, BatchNorm, MHSA-first, relative positional encoding), we use $h = 16$, 1-group GroupNorm, convolution-first ordering, and no positional embeddings. Expanding original ChordFormer's 6-quality head (`triad_limit=6`, 100 logits, collapsing `5/1` into `maj`) to `triad_limit=8` maps $\mathbf{h}_t \in \mathbb{R}^{256}$ to $\hat{\mathbf{z}}_t = [\hat{\mathbf{z}}^{(97)}_{0,t} \mid \hat{\mathbf{z}}^{(13)}_{1,t} \mid \hat{\mathbf{z}}^{(4)}_{2,t} \mid \hat{\mathbf{z}}^{(4)}_{3,t} \mid \hat{\mathbf{z}}^{(3)}_{4,t} \mid \hat{\mathbf{z}}^{(3)}_{5,t}] \in \mathbb{R}^{124}$ (Section 3.1) [3, 9], totaling $|\Theta| = 6{,}123{,}644$ parameters.

During inference (`extractors/xhmm_ismir.py`), posteriors $\mathbf{p}_{h,t} = \text{Softmax}(\hat{\mathbf{z}}_{h,t})$ are decoded via linear-chain CRF Viterbi smoothing over Harte templates [9, 12]:

$$S^*(t) = \arg\max_c \Big\{ \max_{c'} \left(V(t-1, c') - \tau_{\text{crf}} \cdot \mathbb{I}(c' \neq c)\right) + \log P(c \mid \mathbf{x}_t) \Big\}, \tag{5}$$

with $\tau_{\text{crf}} = 2.0$ (multi-head) and $\tau_{\text{crf}} = 30.0$ (MIREX). The CRF suppresses spurious frame-level fluctuations while preserving true harmonic boundaries.

4.2 Datasets and Cross-Validation Protocol

Our experimental workflow employs two corpora: (1) **Synthetic Pretraining Corpus**, used exclusively for pretraining and curriculum replay, comprising 5,000 procedural random-mix tracks from Section 3 (460.3 h, 21.8 GB CQT HDF5; 4,500 train and 500 val tracks, seed `42`); and (2) **Real Benchmark Corpus**, the 1,217-song collection spanning 731 McGill Billboard, 287 USPop2002/MARL, and 199 Isophonics (Beatles/Queen) tracks. The final test set is drawn exclusively from this 1,217 real-song collection (no synthetic or test audio is used in training); however, because our single pretraining corpus extracts symbolic transitions from full ChoCo Billboard/Isophonics tables and calibrates quotas against the 1,217-song class prior (Figure 1), progression statistics from test-fold songs are not isolated from the generator. Under 5-fold cross-validation (seed `20190326`), splits `s0`–`s4` partition the 1,217 real tracks ($730/731$ train, 243 val, $243/244$ test tracks/fold; 11,436,357 frames, 73.7 h), where Fold 0 `maj` (4,465,981 frames) and `min` (1,586,913) outnumber `sus4` (30.5 : 1), `dim` (133.2 : 1), `sus2` (217.5 : 1), `aug` (292.4 : 1), `#9` (221.8 : 1), `b9` (705.2 : 1), `#11` (1,113.1 : 1), and `b13` (1,159.3 : 1).

4.3 Training Regimens: Baseline vs. Sim-to-Real Curriculum

We compare two regimens under FP16 mixed precision and AdamW ($\beta_1 = 0.9, \beta_2 = 0.999$) on $T = 1{,}000$-frame chunks (23.22 s): (1) **Control Baseline** (`control-benchmark`), trained from scratch on the 1,217 real songs for 24 epochs/fold ($B = 36$ or 48, uniform LR η over $\{10^{-3} : 12, 10^{-4} : 8, 10^{-5} : 4\}$ epochs, $[-5, +6]$ pitch shifting); and (2) **Sim-to-Real Curriculum** (`finetuned-benchmark`), pretraining in Stage 1 (`pretrain_synthetic.py`) on 5,000 synthetic tracks for 7 epochs (10,500 mini-batches, $B = 36$, LR $\{10^{-3} : 4, 10^{-4} : 2, 10^{-5} : 1\}$) with unweighted cross-entropy, pitch shifting, and SpecAugment [27] (2 frequency masks ≤ 24 bins, 2 time masks ≤ 60 frames at -80.0 dB), then fine-tuning in Stage 2 (`finetune_real_fold.py`) on each real fold for 24 epochs ($B = 36$, LR $\{10^{-4} : 12, 10^{-5} : 8, 10^{-6} : 4\}$).

Real training minimizes $\mathcal{L}_{\text{total}} = \sum_{h=0}^{5} \lambda_h \mathcal{L}_h(\hat{\mathbf{z}}_h, \mathbf{y}_h; \mathbf{w}_h)$ ($\boldsymbol{\lambda} = [1, 1, 1, 1, 1, 1]$) with inverse-frequency class weights [3, 9]:

$$w_{h,c} = \min\left(\left(\frac{\max_k N_{h,k}}{\max(N_{h,c}, 1.0)}\right)^{\alpha}, C_{\max}\right), \tag{6}$$

using $\alpha = 0.50$, $C_{\max} = 10.0$, and clamps ($w_{2,\text{b7}} \approx 3.602$, $w_{3,9} = 6.000$, $w_{\text{sus4}} \geq 6.800$, $w_5 \geq 6.000$, $w_{\text{min}} \geq 2.000$). Stage 2 adds: (1) **Inverted Discriminative LR** ($\mathbf{s}_{\text{lr}} = [0.2, 0.5, 1.0]$), scaling bottom blocks (Layers 0–1) by 0.2η, top blocks (Layers 2–3) by 0.5η, and output heads by 1.0η; and (2) **Floor Curriculum Replay**, in-

Table 1. 5-fold macro recall and F_1 (%) across the 6 structural classification heads and isolated triad quality subsets (mean ± std across splits s_0–s_4). Control: 1,217-song real baseline; Experimental: 5,000-song synthetic pretraining + real finetuning (Δ = Exp. − Control).

Structural Output Head	Control Rec. (%)	Exp. Rec. (%)	Δ Rec.	Control F_1 (%)	Exp. F_1 (%)	Δ F_1
Primary 6-Head Multi-Task Decomposition:						
Root + Triad (97 classes)	41.21 ± 3.12	**42.66 ± 2.59**	+1.45	36.57 ± 2.07	**36.70 ± 1.68**	+0.13
Bass Pitch (13 classes)	**84.81 ± 1.73**	82.93 ± 1.51	-1.88	**84.72 ± 1.77**	83.01 ± 1.57	-1.70
7th Extension (4 classes)	**62.69 ± 4.27**	62.55 ± 3.14	-0.14	**57.67 ± 2.73**	56.99 ± 1.48	-0.67
9th Extension (4 classes)	35.33 ± 2.68	**43.58 ± 7.54**	+8.25	34.98 ± 2.96	**39.01 ± 6.25**	+4.03
11th Extension (3 classes)	38.74 ± 7.85	**41.71 ± 8.03**	+2.97	37.35 ± 3.60	**42.14 ± 7.04**	+4.79
13th Extension (3 classes)	**35.36 ± 3.74**	34.26 ± 1.95	-1.10	**36.09 ± 4.82**	34.47 ± 2.47	-1.62
Macro Average (6 Heads)	49.69 ± 2.10	**51.28 ± 1.94**	+1.59	47.90 ± 1.69	**48.72 ± 1.86**	+0.82
Targeted Triad Quality Breakdown:						
Triad Quality (8 qualities + N)	50.23 ± 4.53	**53.42 ± 3.96**	+3.19	45.52 ± 3.40	**46.91 ± 2.57**	+1.39
Common Triads (*maj*, *min*)	**84.74 ± 1.97**	83.08 ± 1.44	-1.66	**86.99 ± 1.19**	86.05 ± 0.85	-0.94
Rare Triads (*sus2*, *sus4*, *dim*, *aug*)	41.67 ± 7.30	**48.31 ± 6.47**	+6.64	53.11 ± 8.62	**60.94 ± 6.26**	+7.83

jecting SpecAugmented synthetic chunks at $\rho(\eta) = 0.150$ ($\eta \geq 10^{-4}$), 0.125 ($\eta = 10^{-5}$), and a strict floor $\rho_{\min} = 0.100$ ($\eta \leq 10^{-6}$) against catastrophic forgetting.

5. RESULTS AND DISCUSSION

5.1 Structured Head-Wise Recall and Rare-Class Gains

Table 1 summarizes 5-fold cross-validated performance ($\tau_{\text{crf}} = 2.0$) across the six structural heads over all 11,436,357 test frames. Across the 6-head macro average, our Sim-to-Real curriculum raises recall from $49.69 \pm 2.10\%$ to $51.28 \pm 1.94\%$ (+1.59 pp) and macro F_1 from $47.90 \pm 1.69\%$ to $48.72 \pm 1.86\%$ (+0.82 pp), concentrating gains on sample-starved classes.

On Head 0 (Root + Triad, 97 classes), macro recall increases from $41.21 \pm 3.12\%$ to $42.66 \pm 2.59\%$ (+1.45 pp) and F_1 from $36.57 \pm 2.07\%$ to $36.70 \pm 1.68\%$ (+0.13 pp), driven by non-diatonic structures: overall triad recall rises from $50.23 \pm 4.53\%$ to $53.42 \pm 3.96\%$ (+3.19 pp; F_1 $45.52 \pm 3.40\% \rightarrow 46.91 \pm 2.57\%$). While common triads (`maj`, `min`) shift by −1.66 pp in recall ($84.74 \pm 1.97\% \rightarrow 83.08 \pm 1.44\%$) and −0.94 pp in F_1 ($86.99 \pm 1.19\% \rightarrow 86.05 \pm 0.85\%$), rare triads (`sus2`, `sus4`, `dim`, `aug`) surge from $41.67 \pm 7.30\%$ to $48.31 \pm 6.47\%$ in recall (+6.64 pp, +15.9% relative) and from $53.11 \pm 8.62\%$ to $60.94 \pm 6.26\%$ in F_1 (+7.83 pp, +14.7% relative). Similarly, Head 3 (9th) recall jumps from $35.33 \pm 2.68\%$ to $43.58 \pm 7.54\%$ (+8.25 pp, +23.4% relative; F_1 +4.03 pp to $39.01 \pm 6.25\%$), and Head 4 (11th) rises from $38.74 \pm 7.85\%$ to $41.71 \pm 8.03\%$ (+2.97 pp; F_1 +4.79 pp to $42.14 \pm 7.04\%$).

Per-class metrics (Figure 2) confirm large gains across rare qualities and extensions: `aug` recall rises from 43.37% to 61.02% (+17.65 pp, F_1 $35.40\% \rightarrow 43.15\%$), `sus2` from 17.32% to 30.25% (+12.93 pp, F_1 $17.40\% \rightarrow 19.90\%$), standalone `5` (uncollapsed from `maj`) from 28.76% to 41.91% (+13.15 pp, F_1 $28.49\% \rightarrow 29.50\%$), `dim` from 65.61% to 67.97% (+2.36 pp, F_1 $41.65\% \rightarrow 47.79\%$), `sus4` from 51.19% to 53.72% (+2.53 pp), and `b7` from 70.31% to 80.90% (+10.59 pp; `bb7` recall +0.36 pp, F_1 $23.15\% \rightarrow 25.22\%$). On upper extensions, natural `9` leaps +25.16 pp ($31.64\% \rightarrow 56.80\%$, F_1 $31.06\% \rightarrow 36.19\%$); `#9` rises 79× ($0.07\% \rightarrow 5.53\%$, F_1 $0.10\% \rightarrow 4.97\%$); `b9` rises +6.32 pp ($18.39\% \rightarrow 24.71\%$, F_1 $9.56\% \rightarrow 18.02\%$); `#11` rises +8.87 pp ($14.15\% \rightarrow 23.02\%$, F_1 $13.39\% \rightarrow 20.50\%$); and `b13` jumps 7.8× ($0.84\% \rightarrow 6.58\%$, F_1 $0.77\% \rightarrow 8.43\%$). Conversely, recall regresses on bare-root `1` (−3.61 pp), major `7` (−3.10 pp), and natural `13` (−2.59 pp, $3.35\% \rightarrow 0.76\%$).

Table 2. Weighted Chord Symbol Recall (WCSR, %) across 5-fold cross-validation on 1,217 real songs (s_0–s_4, mean ± std). Control: 1,217-song real baseline; Experimental: synthetic pretraining + real finetuning.

Evaluation Rule	Control (%)	Experimental (%)	Δ (pp)
Root Pitch Only	**83.77 ± 1.88**	82.30 ± 1.58	-1.47
Major / Minor	**79.78 ± 2.38**	77.04 ± 1.71	-2.74
Major / Minor + Inversions	**77.77 ± 2.22**	75.17 ± 1.71	-2.60
Triads (Full Quality)	**75.01 ± 2.72**	72.88 ± 1.97	-2.13
Triads + Inversions	**73.15 ± 2.58**	71.13 ± 1.95	-2.02
Sevenths (Standard)	**66.46 ± 2.46**	62.46 ± 1.73	-4.00
Sevenths + Inversions	**65.12 ± 2.42**	61.18 ± 1.72	-3.94
Tetrads (Extended)	**61.62 ± 2.73**	58.20 ± 1.74	-3.42
Tetrads + Inversions	**60.38 ± 2.69**	57.02 ± 1.75	-3.36
MIREX Metric (Overall)	**81.49 ± 2.56**	79.59 ± 1.85	-1.90

Both models share identical inverse-frequency reweighting ($\alpha = 0.50, C_{\max} = 10.0$) in real training, isolating the effect of synthetic pretraining at this operating point; we hypothesize that without diverse synthetic exemplars, amplifying gradients on a handful of real tracks overfits to track-specific timbres rather than learning generalizable minority-class boundaries.

5.2 Standard Frame-Level WCSR and MIREX Evaluation

Table 2 reports duration-weighted chord symbol recall (WCSR) via `mir_eval.chord` [28] ($\tau_{\text{crf}} = 30.0$). Across all ten comparison rules, the Sim-to-Real curriculum scores strictly lower than the real-only baseline: Root drops by 1.47 pp (83.77% → 82.30%), Major/Minor by 2.74 pp (79.78% → 77.04%; −2.60 pp with inversions), Triads by 2.13 pp (75.01% → 72.88%; −2.02 pp with inversions), Sevenths by 4.00 pp (66.46% → 62.46%), Tetrads by 3.42 pp (61.62% → 58.20%), and overall MIREX WCSR by 1.90 pp ($81.49 \pm 2.56\% \rightarrow 79.59 \pm 1.85\%$).

This divergence reflects a fundamental precision-recall and metric-weighting trade-off: WCSR weights accuracy strictly by interval duration ($\sum_i d_i \mathbb{I}(\hat{C}_i == C_i) / \sum_i d_i$), biasing it toward unextended `maj` (44.93% of real duration) and `min` (11.95%; 65.75%/22.16% on Head 0 with extensions). Under Harte recombination, candidate

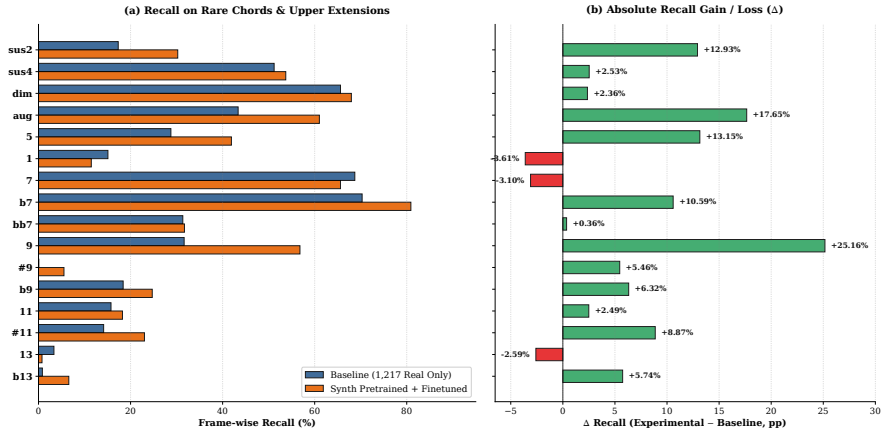


Figure 2. Per-class evaluation on rare triads, power chords, 7th extensions, and upper/altered extensions across 11,436,357 frames of the 1,217-song real benchmark: (a) absolute recall comparing Control Baseline (real-only) versus Experimental Curriculum (synthetic pretraining + real finetuning); (b) diverging recall gain (ΔRecall), highlighting gains across `9` (+25.16 pp), `aug` (+17.65 pp), `5` (+13.15 pp), `sus2` (+12.93 pp), `b7` (+10.59 pp), `#11` (+8.87 pp), `b9` (+6.32 pp), `b13` (+5.74 pp), and `#9` (+5.46 pp).

log-probabilities sum across all six heads ($\log P_k(t) = \log p_{\text{triad}} + \mathbb{I}(\text{use_bass}) \log p_{\text{bass}} + \sum_e \log p_e$). Expanding minority decision boundaries in $\hat{\mathbf{z}}_{h,t}$ to capture rare qualities and extensions inevitably introduces false-positive activations on acoustically ambiguous frames of dominant triads (e.g., guitar overtones or vocal passing tones triggering `sus4` or `b7`, where per-class F_1 shifts by -1.79 pp and -0.89 pp despite $+2.53$ pp and $+10.59$ pp recall gains). Because rare classes occupy minuscule duration slices (`sus2` 0.34%, `aug` 0.21%, `b9` 0.13%, `b13` 0.08%), true-positive detections on the harmonic tail yield negligible WCSR credit, whereas any false-positive excursion on a majority triad immediately penalizes duration overlap.

Stage 2 adaptation bounds this trade-off. Updating linear heads at 1.0η while anchoring bottom Conformer blocks (0.2η) and maintaining the $\rho_{\min} = 0.10$ synthetic replay floor recalibrates output logit priors to studio mixes without erasing minority features. Moreover, synthetic pretraining reduces 5-fold standard deviation across all ten rules in Table 2 (MIREX $\pm 2.56\% \rightarrow \pm 1.85\%$; MajMin $\pm 2.38\% \rightarrow \pm 1.71\%$; Tetrads $\pm 2.73\% \rightarrow \pm 1.74\%$), confirming cross-fold stability.

5.3 Sim-to-Real Acoustic Domain Gap and Extension Masking

Bridging SoundFont synthesis to studio recordings requires overcoming static synthesis timbres and polyphonic overtone masking. Because encoders can memorize deterministic SoundFont attacks and envelopes, our random-mix pipeline applies Bernoulli stem dropout ($P = 0.70$ on lead/drums), gain perturbations ($\Delta G \in [-4, +4]$ dB), overdrive, split-band room convolution, and detuning ($\Delta f_{\text{cents}} \sim U(-25, +25)$ cents), forcing the $K = 31$ convolutions and attention heads to learn timbre-invariant intervals across detuned CQT bins.

Furthermore, we hypothesize that physical register and overtone alignment help explain the heterogeneous per-class trends in Figure 2: (1) *Low-Register Triads*: Rare triads (`sus2`, `aug`, `dim`) alter third/fifth degrees in C_1–C_4, where 36 bins/octave CQT filterbanks resolve fundamentals cleanly (+2.36 pp to +17.65 pp recall), whereas bare-root `1` (−3.61 pp) lacks tertian intervals to distinguish it from bass/melody octaves; (2) *Natural Extensions* (`7`, `11`, `13`): In C_4–B_6, we hypothesize that upper tones compete with vocal formants, guitars, and cymbals—natural `11` (17 semitones) clashes against a major 3rd (+2.49 pp), while major `7` (−3.10 pp) and natural `13` (−2.59 pp, $3.35\% \rightarrow 0.76\%$) have partials that align closely with root and fifth overtones; and (3) *Altered Extensions* (`#9`, `b9`, `#11`, `b13`): Chromatic alterations diverge from the diatonic overtone series of the triad, which may produce more isolated CQT peaks (79× on `#9`, 7.8× on `b13`).

6. CONCLUSION

We presented a Sim-to-Real pretraining framework for large-vocabulary ACR that synthesizes 5,000 multitracks (460.3 h) via empirical transition sampling, rare-chord quota scheduling, physically constrained voicings, and random-mix augmentation. Pretraining ChordFormer on synthetic audio and finetuning with inverted discriminative learning rates ($\mathbf{s}_{\text{lr}} = [0.2, 0.5, 1.0]$) and floor curriculum replay ($\rho_{\min} = 0.10$) yields substantial recall gains over the real-only baseline on rare triads (+6.64 pp), natural 9ths (+25.16 pp), and alterations (79× on `#9`, 7.8× on `b13`) while reducing 5-fold WCSR variance (code/models: `https://github.com/shojha24/probabilistic-composition-generator`, `https://github.com/shojha24/ChordFormer-Artificial-Dataset-Benchmarking`). Limitations of these preliminary results include testing a single quota operating point and extracting generator transitions/priors from full ChoCo Billboard/Isophonics tables rather than per-fold training splits; future work will enforce fold-isolated extraction, ablate loss reweighting, tune quotas for `1`/`7`/`13`, and integrate DDSP.

7. REFERENCES

[1] J. Pauwels, K. O'Hanlon, E. Gómez, and M. B. Sandler, "20 years of automatic chord recognition from audio," in *Proceedings of the 20th International Society for Music Information Retrieval Conference (ISMIR)*, Delft, The Netherlands, 2019, pp. 54–63.

[2] F. Korzeniowski and G. Widmer, "Feature learning for chord recognition: The deep chroma extractor," in *Proceedings of the 17th International Society for Music Information Retrieval Conference (ISMIR)*, New York City, USA, 2016, pp. 37–43.

[3] J. Jiang, K. Chen, W. Li, and G. Xia, "Large-vocabulary chord transcription via chord structure decomposition," in *Proceedings of the 20th International Society for Music Information Retrieval Conference (ISMIR)*, Delft, The Netherlands, 2019, pp. 644–651.

[4] J. Park, K. Choi, S. Jeon, D. Kim, and J. Park, "A bi-directional Transformer for musical chord recognition," in *Proceedings of the 20th International Society for Music Information Retrieval Conference (ISMIR)*, Delft, The Netherlands, 2019, pp. 620–627.

[5] A. Gulati, J. Qin, C.-C. Chiu, N. Parmar, Y. Zhang, J. Yu, W. Han, S. Wang, Z. Zhang, Y. Wu, and R. Pang, "Conformer: Convolution-augmented transformer for speech recognition," in *Proceedings of Interspeech 2020*, 2020, pp. 5036–5040.

[6] E. J. Humphrey and J. P. Bello, "Four timely insights on automatic chord estimation," in *Proceedings of the 16th International Society for Music Information Retrieval Conference (ISMIR)*, Málaga, Spain, 2015, pp. 673–679.

[7] H. V. Koops, W. B. de Haas, J. A. Burgoyne, J. Bransen, A. Kent-Muller, and A. Volk, "Annotator subjectivity in harmony annotations of popular music," *Journal of New Music Research*, vol. 48, no. 3, pp. 232–252, 2019.

[8] B. McFee and J. P. Bello, "Structured training for large-vocabulary chord recognition," in *Proceedings of the 18th International Society for Music Information Retrieval Conference (ISMIR)*, Suzhou, China, 2017, pp. 188–194.

[9] M. W. Akram, S. Dettori, V. Colla, and G. C. Buttazzo, "ChordFormer: A conformer-based architecture for large-vocabulary audio chord recognition," *IEEE Transactions on Audio, Speech, and Language Processing*, vol. 34, pp. 581–595, 2026.

[10] C. Li, J. Jiang, Y. Li, and L. Tian, "Large-vocabulary chord recognition based on contrastive learning and noisy student," *IEEE Transactions on Consumer Electronics*, vol. 71, no. 2, pp. 3695–3706, 2025.

[11] J. de Berardinis, A. Meroño-Peñuela, A. Poltronieri, and V. Presutti, "ChoCo: a chord corpus and a data transformation workflow for musical harmony knowledge graphs," *Scientific Data*, vol. 10, no. 1, p. 641, 2023.

[12] C. Harte, M. B. Sandler, S. A. Abdallah, and E. Gómez, "Symbolic representation of musical chords: A proposed syntax for text annotations," in *Proceedings of the 6th International Society for Music Information Retrieval Conference (ISMIR)*, London, UK, 2005, pp. 66–71.

[13] L. Kim and J. Park, "An event-based sequence modeling approach to recognizing non-triad chords with oversegmentation minimization," in *Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2026, pp. 15 277–15 281.

[14] E. Manilow, G. Wichern, P. Seetharaman, and J. Le Roux, "Cutting music source separation some Slakh: A dataset to study the impact of training data quality and quantity," in *Proceedings of the IEEE Workshop on Applications of Signal Processing to Audio and Acoustics (WASPAA)*, New Paltz, NY, USA, 2019, pp. 45–49.

[15] F. Ostermann, I. Vatolkin, and M. Ebeling, "AAM: a dataset of Artificial Audio Multitracks for diverse music information retrieval tasks," *EURASIP Journal on Audio, Speech, and Music Processing*, vol. 2023, no. 1, p. 13, 2023.

[16] M. Majchrzak and J. Mańdziuk, "Training chord recognition models on artificially generated audio," *Neural Computing and Applications*, vol. 38, p. 240, 2026.

[17] J. A. Burgoyne, J. Wild, and I. Fujinaga, "An expert ground-truth set for audio chord recognition and music analysis," in *Proceedings of the 12th International Society for Music Information Retrieval Conference (ISMIR)*, Miami, Florida, USA, 2011, pp. 633–638.

[18] M. Mauch, C. Cannam, M. Davies, S. Dixon, C. Harte, S. Kolozali, D. Tidhar, and M. Sandler, "OMRAS2 metadata project 2009," in *Proceedings of the 10th International Society for Music Information Retrieval Conference (ISMIR) Late-Breaking Demos*, Kobe, Japan, 2009.

[19] M. Pfleiderer, K. Frieler, J. Abeßer, W.-G. Zaddach, and B. Burkhart, Eds., *Inside the Jazzomat: New Perspectives for Jazz Research*. Mainz, Germany: Schott Campus, 2017.

[20] D. Tymoczko, *A Geometry of Music: Harmony and Counterpoint in the Extended Common Practice*, ser. Oxford Studies in Music Theory. New York, NY, USA: Oxford University Press, 2011.

[21] C.-R. Nadar, J. Abeßer, and S. Grollmisch, "Towards CNN-based acoustic modeling of seventh chords for automatic chord recognition," in *Proceedings of the*

16th Sound and Music Computing Conference (SMC), Málaga, Spain, 2019, pp. 48–54.

[22] D. Tymoczko, "The geometry of musical chords," *Science*, vol. 313, no. 5783, pp. 72–74, 2006.

[23] M. Eitel, N. Ruth, P. Harrison, K. Frieler, and D. Müllensiefen, "Perception of chord sequences modeled with prediction by partial matching, voice-leading distance, and spectral pitch-class similarity: A new approach for testing individual differences in harmony perception," *Music & Science*, vol. 7, pp. 1–16, 2024.

[24] R. Cohn, "Introduction to neo-riemannian theory: A survey and a historical perspective," *Journal of Music Theory*, vol. 42, no. 2, pp. 167–180, 1998.

[25] J. Douthett and P. Steinbach, "Parsimonious graphs: A study in parsimony, contextual transformations, and modes of limited transposition," *Journal of Music Theory*, vol. 42, no. 2, pp. 241–263, 1998.

[26] J. C. Brown, "Calculation of a constant Q spectral transform," *The Journal of the Acoustical Society of America*, vol. 89, no. 1, pp. 425–434, 1991.

[27] D. S. Park, W. Chan, Y. Zhang, C.-C. Chiu, B. Zoph, E. D. Cubuk, and Q. V. Le, "SpecAugment: A simple data augmentation method for automatic speech recognition," in *Proceedings of Interspeech 2019*, 2019, pp. 2613–2617.

[28] C. Raffel, B. McFee, E. J. Humphrey, J. Salamon, O. Nieto, D. Liang, and D. P. W. Ellis, "mir_eval: A transparent implementation of common MIR metrics," in *Proceedings of the 15th International Society for Music Information Retrieval Conference (ISMIR)*, Taipei, Taiwan, 2014, pp. 367–372.